

## BREAST CANCER DIAGNOSIS AND PREDICTION USING THE WISCONSIN BREAST CANCER DIAGNOSTIC DATASET

### Library Imports

```
In [2]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.metrics import confusion_matrix
from sklearn.metrics import roc_auc_score, roc_curve
from sklearn.metrics import classification_report
from sklearn.metrics import accuracy_score
from sklearn.linear_model import LogisticRegression
lr = LogisticRegression()
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
Rf = RandomForestClassifier()
from sklearn.neighbors import KNeighborsClassifier
knn = KNeighborsClassifier(n_neighbors=12)
from sklearn.svm import SVC
svm = SVC(kernel='linear', random_state = 10)
```

### Dataset Upload

```
In [4]: df = pd.read_csv('C:/Users/USER/Desktop/breast cancer wiscosin.csv')
```

```
In [5]: # Display of the first 5 rows in the dataset
df.head(5)
```

```
Out[5]:
```

	id	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean	smoothness_mea
0	842302	M	17.99	10.38	122.80	1001.0	0.1184
1	842517	M	20.57	17.77	132.90	1326.0	0.0847
2	84300903	M	19.69	21.25	130.00	1203.0	0.1096
3	84348301	M	11.42	20.38	77.58	386.1	0.1425
4	84358402	M	20.29	14.34	135.10	1297.0	0.1003

5 rows × 32 columns

```
In [6]: #Overview of the Information in the Dataset
df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 569 entries, 0 to 568
Data columns (total 32 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   id                                     569 non-null    int64
1   diagnosis                             569 non-null    object
2   radius_mean                           569 non-null    float64
3   texture_mean                           569 non-null    float64
4   perimeter_mean                         569 non-null    float64
5   area_mean                             569 non-null    float64
6   smoothness_mean                       569 non-null    float64
7   compactness_mean                      569 non-null    float64
8   concavity_mean                        569 non-null    float64
9   concave points_mean                   569 non-null    float64
10  symmetry_mean                         569 non-null    float64
11  fractal_dimension_mean                569 non-null    float64
12  radius_se                             569 non-null    float64
13  texture_se                             569 non-null    float64
14  perimeter_se                           569 non-null    float64
15  area_se                               569 non-null    float64
16  smoothness_se                         569 non-null    float64
17  compactness_se                        569 non-null    float64
18  concavity_se                          569 non-null    float64
19  concave points_se                     569 non-null    float64
20  symmetry_se                           569 non-null    float64
21  fractal_dimension_se                  569 non-null    float64
22  radius_worst                          569 non-null    float64
23  texture_worst                         569 non-null    float64
24  perimeter_worst                       569 non-null    float64
25  area_worst                            569 non-null    float64
26  smoothness_worst                      569 non-null    float64
27  compactness_worst                     569 non-null    float64
28  concavity_worst                       569 non-null    float64
29  concave points_worst                  569 non-null    float64
30  symmetry_worst                        569 non-null    float64
31  fractal_dimension_worst                569 non-null    float64
dtypes: float64(30), int64(1), object(1)
memory usage: 142.4+ KB

```

### Data Preprocessing

```

In [7]: #Checking For Missing Values
df.isnull().sum().sort_values(ascending = False)

```

```
Out[7]: id 0
diagnosis 0
symmetry_worst 0
concave points_worst 0
concavity_worst 0
compactness_worst 0
smoothness_worst 0
area_worst 0
perimeter_worst 0
texture_worst 0
radius_worst 0
fractal_dimension_se 0
symmetry_se 0
concave points_se 0
concavity_se 0
compactness_se 0
smoothness_se 0
area_se 0
perimeter_se 0
texture_se 0
radius_se 0
fractal_dimension_mean 0
symmetry_mean 0
concave points_mean 0
concavity_mean 0
compactness_mean 0
smoothness_mean 0
area_mean 0
perimeter_mean 0
texture_mean 0
radius_mean 0
fractal_dimension_worst 0
dtype: int64
```

```
In [8]: #Check for Duplicates
df = df.drop_duplicates()
df.shape
#No duplicates in the Dataset
```

```
Out[8]: (569, 32)
```

```
In [9]: #Dropping the ID Column
df = df.drop(['id'], axis = 1)
```

```
In [12]: df.diagnosis.replace({'B':'Benign','M':'Malignant'}, inplace = True)
```

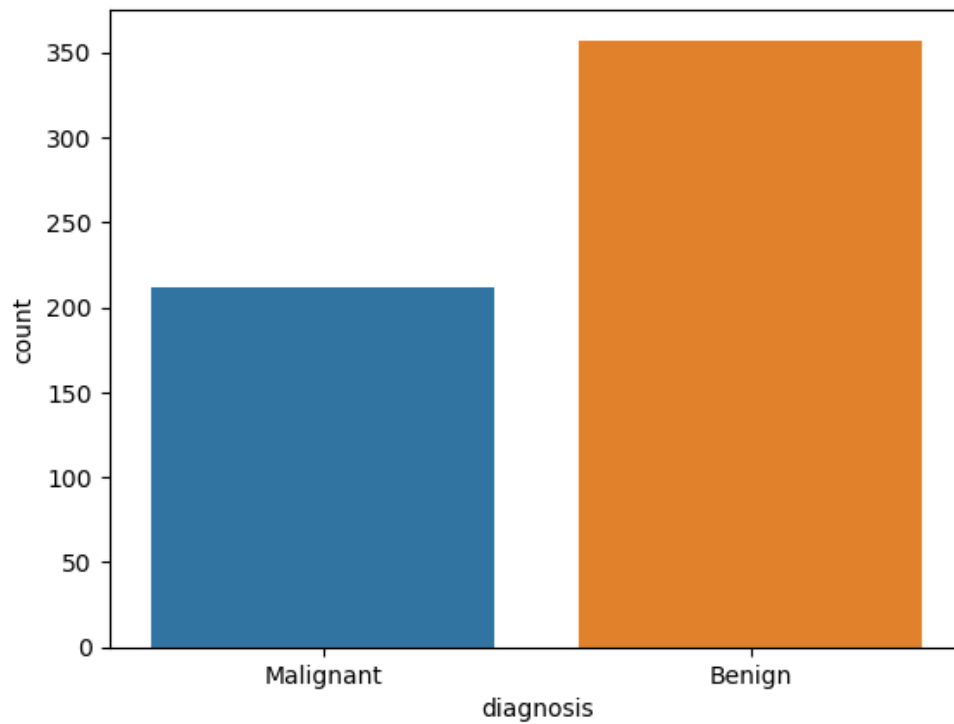
```
In [13]: df.diagnosis.value_counts()
```

```
Out[13]: Benign      357
Malignant    212
Name: diagnosis, dtype: int64
```

```
In [14]: sns.countplot(x='diagnosis', data = df)
plt.title('WISCONSIN BREAST CANCER DIAGNOSTIC DATASET DISPLAYING THE DIAGNOSIS')
```

```
Out[14]: Text(0.5, 1.0, 'WISCONSIN BREAST CANCER DIAGNOSTIC DATASET DISPLAYING THE DIAGNOSIS')
```

## WISCONSIN BREAST CANCER DIAGNOSTIC DATASET DISPLAYING THE DIAGNOSIS




```
In [28]: #Converting Categorical Data to Ordinal  
df.diagnosis.replace({'Benign':0, 'Malignant':1}, inplace = True)
```

## FEATURE SELECTION WITH CORRELATION MATRIX

```
In [16]: correlation_matrix = df.corr()
```

```
In [17]: plt.figure(figsize=(18, 18))  
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', linewidths=3)
```

```
Out[17]: <AxesSubplot:>
```



```
In [18]: threshold = 0.8
correlated_pairs = {}

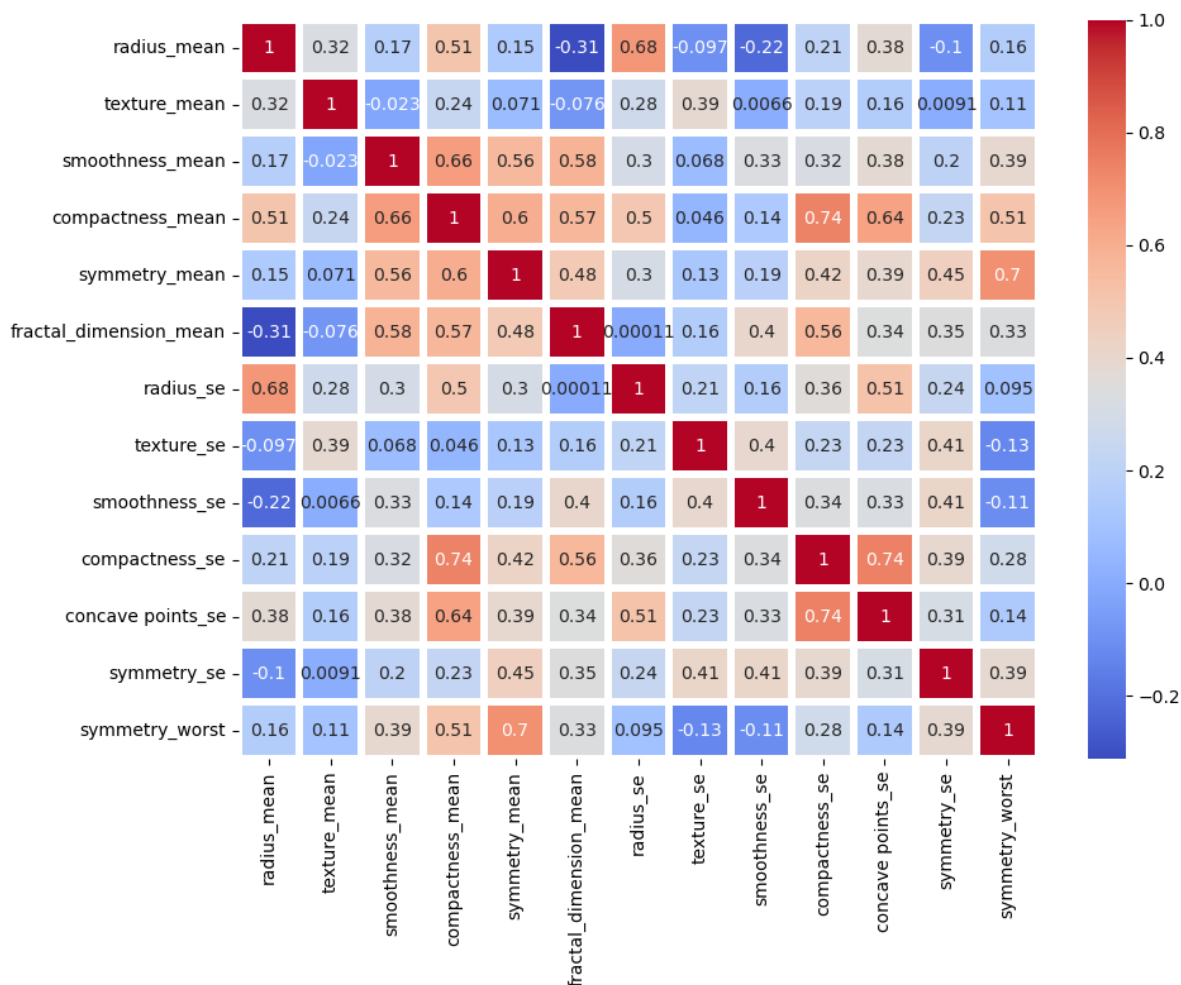
for i in range(len(correlation_matrix.columns)):
    for j in range(i):
        if abs(correlation_matrix.iloc[i, j]) > threshold:
            colname = correlation_matrix.columns[i]
            correlated_pairs[colname] = correlation_matrix.columns[j]
```

```
In [19]: #Removal of Select Features With High Correlation
features_to_remove = set(correlated_pairs.keys())
df_reduced = df.drop(columns=features_to_remove)
```

```
In [20]: #Dataset with reduced features
df_reduced.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 569 entries, 0 to 568
Data columns (total 14 columns):
#   Column                                Non-Null Count  Dtype
---  ---                                -
0   diagnosis                             569 non-null    object
1   radius_mean                           569 non-null    float64
2   texture_mean                           569 non-null    float64
3   smoothness_mean                       569 non-null    float64
4   compactness_mean                      569 non-null    float64
5   symmetry_mean                         569 non-null    float64
6   fractal_dimension_mean               569 non-null    float64
7   radius_se                             569 non-null    float64
8   texture_se                             569 non-null    float64
9   smoothness_se                        569 non-null    float64
10  compactness_se                       569 non-null    float64
11  concave points_se                    569 non-null    float64
12  symmetry_se                          569 non-null    float64
13  symmetry_worst                      569 non-null    float64
dtypes: float64(13), object(1)
memory usage: 66.7+ KB
```

```
In [21]: #Heatmap displaying the current features and their correlation
cm = df_reduced.corr()
plt.figure(figsize=(10, 8))
sns.heatmap(cm, annot=True, cmap='coolwarm', linewidths=3)
plt.tight_layout()
```



TRAIN - TEST SPLIT (75:25)

```
In [22]: df_reduced = df_reduced.drop(['diagnosis'], axis = 1)
```

```
In [29]: X = df_reduced
y = df['diagnosis']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25, random_st
```

### Model Implementation

```
In [30]: # Defining a function to plot the Receiver Operating Characteristic (ROC) curve
def plot_roc_curve(y_true, y_scores):
    fpr, tpr, _ = roc_curve(y_true, y_scores)
    auc = roc_auc_score(y_true, y_scores)
    plt.figure()
    plt.plot(fpr, tpr, label='ROC curve (AUC = {:.2f})'.format(auc))
    plt.xlim([0.0, 1.0])
    plt.ylim([0.0, 1.05])
    plt.xlabel('False Positive Rate (FPR)')
    plt.ylabel('True Positive Rate (TPR)')
    plt.title('Receiver Operating Characteristic (ROC) Curve')
    plt.legend(loc='lower right')
    plt.show()
```


### Logistic Regression Model

```
In [31]: lr = LogisticRegression()
lr.fit(X_train, y_train)
predictions = lr.predict(X_test)
print(confusion_matrix(y_test, predictions))
print(classification_report(y_test, predictions))
print("ROC_AUC Score : ", '{0:.2%}'.format(roc_auc_score(y_test, predictions)))
plot_roc_curve(y_test, predictions)
```

```
[[90  0]
 [ 4 49]]
```

	precision	recall	f1-score	support
0	0.96	1.00	0.98	90
1	1.00	0.92	0.96	53
accuracy			0.97	143
macro avg	0.98	0.96	0.97	143
weighted avg	0.97	0.97	0.97	143

ROC\_AUC Score : 96.23%




```
In [32]: Rf.fit(X_train, y_train)
predictions = Rf.predict(X_test)
print(confusion_matrix(y_test, predictions))
print(classification_report(y_test, predictions))
print("ROC_AUC Score : ", '{0:.2%}'.format(roc_auc_score(y_test, predictions)))
plot_roc_curve(y_test, predictions)
```

```
[[88  2]
 [ 3 50]]
```

	precision	recall	f1-score	support
0	0.97	0.98	0.97	90
1	0.96	0.94	0.95	53
accuracy			0.97	143
macro avg	0.96	0.96	0.96	143
weighted avg	0.96	0.97	0.96	143

ROC\_AUC Score : 96.06%






```
In [33]: svm.fit(X_train, y_train)
predictions = svm.predict(X_test)
print(confusion_matrix(y_test, predictions))
print(classification_report(y_test, predictions))
print("ROC_AUC Score : ", '{0:.2%}'.format(roc_auc_score(y_test, predictions)))
plot_roc_curve(y_test, predictions)
```

```
[[89  1]
 [ 2 51]]
```

	precision	recall	f1-score	support
0	0.98	0.99	0.98	90
1	0.98	0.96	0.97	53
accuracy			0.98	143
macro avg	0.98	0.98	0.98	143
weighted avg	0.98	0.98	0.98	143

ROC\_AUC Score : 97.56%



```
In [34]: Dt_model = DecisionTreeClassifier(max_depth = 7)
Dt_model.fit(X_train, y_train)
predictions = Dt_model.predict(X_test)
print(confusion_matrix(y_test, predictions))
print(classification_report(y_test, predictions))
print("ROC_AUC Score : ", '{0:.2%}'.format(roc_auc_score(y_test, predictions)))
plot_roc_curve(y_test, predictions)
```

```
[[86  4]
 [ 4 49]]
```

	precision	recall	f1-score	support
0	0.96	0.96	0.96	90
1	0.92	0.92	0.92	53
accuracy			0.94	143
macro avg	0.94	0.94	0.94	143
weighted avg	0.94	0.94	0.94	143

ROC\_AUC Score : 94.00%

